



# Facebook

22z433

Santhosh TK

23z431

Dwarkesh

23z432

Kapil arif K

23z434

Praveen G

23z435

Sudharsan S

23z436

VasanthKumaar SB



# *Introduction to facebook*

**23z434**

**PRAVEEN  
G**




# ***INTRODUCTION:***

---

- ❑ Facebook is one of the largest social networking platforms, connecting over 3 billion users worldwide.
- ❑ Every second, Facebook processes a massive volume of data—from posts, messages, comments, to photos and videos.
- ❑ To handle this enormous scale efficiently, Facebook relies on distributed computing, which helps manage large data, ensure fast performance, and provide reliable services to users around the globe.

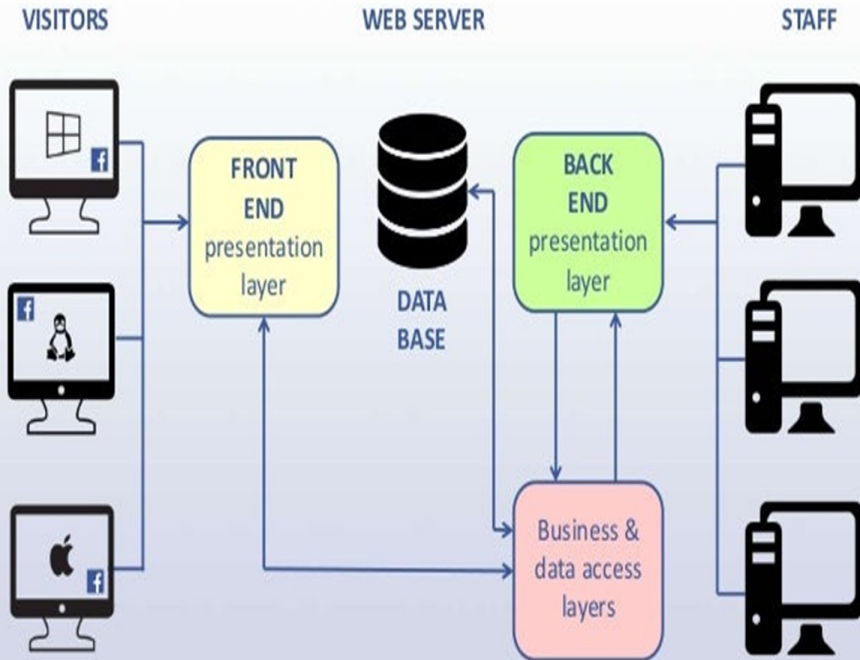
# Why Facebook Needs Distributed Computing:

- ❑ **Massive Scale:** With billions of users, Facebook needs systems that can process huge amounts of data simultaneously.
- ❑ **Load Balancers:** Automatically redirect traffic if one server goes down, so users don't experience downtime.
- ❑ **Backup Systems:** Regular backups ensure that data can be restored quickly in case of a failure.
- ❑ **Monitoring Tools:** Systems are constantly monitored to detect and fix issues before they affect users.
- ❑ **Distributed Systems:** Spreading data and services across multiple locations reduces the risk of a single point of failure.



# ARCHITECTURE OVERVIEW

## Front end & Back end



Facebook's **frontend** is built using the **LAMP stack**.



## BACKEND :TECHNOLOGY STACKS

**Programming Languages:** C+  
+,Python,Java,Erlang

## LAMP stack, which includes:

---



- **Linux** for server operations.
- **Apache** as the web server to handle incoming requests.
- **MySQL** for managing user data.
- **PHP** to create dynamic and interactive web content.

## Why LAMP?

- It's **scalable** and can handle large traffic volumes.
- Being **open-source**, it allows flexibility and cost efficiency.
- Reliable for serving **dynamic content** quickly.




# **Introduction of Facebook's Infrastructure**

**K . KAPIL  
ARIF  
23Z432**

---

Facebook operates at an unprecedented scale, requiring highly efficient, distributed systems to handle massive amounts of data and billions of user interactions daily. Traditional database and storage architectures are not sufficient for such a large-scale system. Instead, Facebook has developed specialized solutions to handle different types of workloads efficiently.



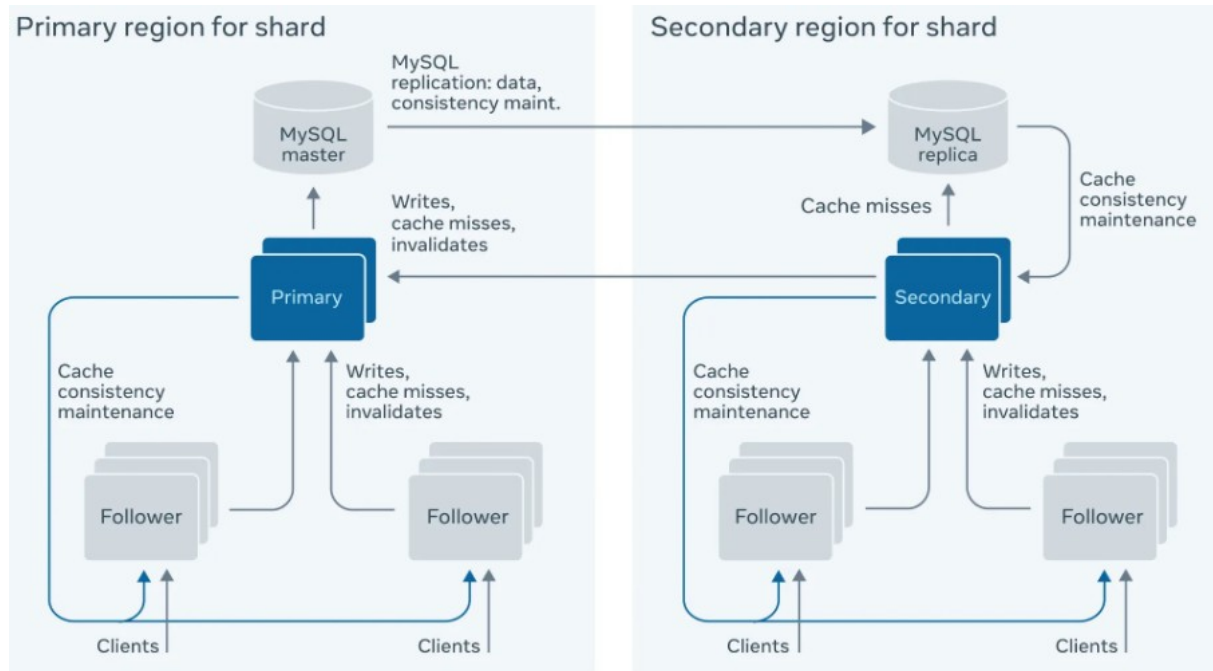


# TAO (The Associations and Objects System)

TAO is Facebook's distributed graph storage system, designed to store and manage relationships between users, posts, and interactions efficiently.

Traditional relational databases struggle to handle such large-scale relationship queries efficiently.

TAO enables fast reads and writes, allowing users to retrieve and update their connections in real-time.



---

## **Primary Region for Shard:**

- Clients interact with follower nodes that provide cached data.
- If the cache is missed, the request is forwarded to the Primary cache.
- If the data is still not found, a query is made to the MySQL master database.
- The response is cached and returned to the clients.
- Any writes, cache misses, or invalidations go through the Primary node to maintain data consistency.

## **Secondary Region for Shard:**

- Similar to the Primary Region, but it uses a MySQL replica instead of the master database.
- Handles cache consistency maintenance by synchronizing with the Primary Region.
- Cache misses are forwarded to the Secondary cache and, if necessary, to the MySQL replica.

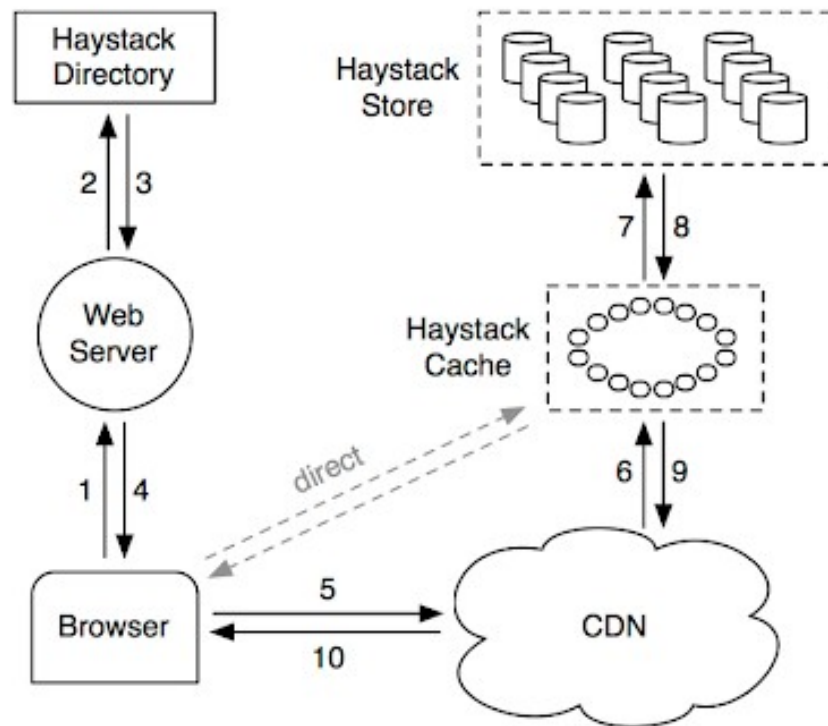
# Haystack

Haystack is Facebook's high-performance photo storage system, designed to efficiently store and retrieve billions of small image files.

Facebook users upload over 350 million photos every day.

Storing images in traditional file systems results in slow access due to frequent disk operations.

Haystack minimizes disk access and retrieves images in a fraction of a second.



---

## WORKING

1. Browser requests a photo → The request goes to the web server.
2. Web server queries the Haystack Directory to locate the requested photo.
3. Haystack Directory responds with the location of the photo.
4. Web server checks if the photo is available in cache/CDN(Content Delivery Network).
5. If the photo is already in the CDN, it is delivered instantly to the browser .
6. If not, the CDN requests the photo from Haystack Cache .
7. If the photo is in Haystack Cache, it is delivered to the CDN and then sent to the browser.
8. The CDN delivers the photo to the browser

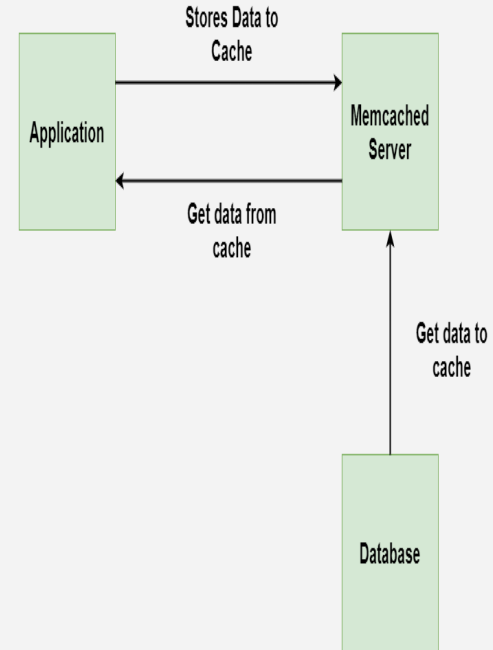
# Memcached

Memcached is Facebook's distributed caching system that stores frequently accessed data in memory to reduce database load and improve response times.

Facebook serves billions of requests per second.

If every request queried the database, it would cause slowdowns and high server costs.

Memcached stores frequently accessed data in RAM, making retrieval almost instantaneous



---

## 1. **Application Requests Data:**

- When an application needs data, it first checks Memcached Server (cache) to see if the data is available.

## 2. **Cache Hit (Data Found in Memcached):**

- If the data exists in Memcached, it is retrieved directly from the cache, avoiding a database query.
- This speeds up the application's response time.

## 3. **Cache Miss (Data Not Found in Memcached):**

- If the requested data is not in Memcached, the application queries the Database.

## 4. **Data Retrieval & Caching:**

- The Database fetches the requested data.
- The Memcached Server stores this data for future requests.
- The Application receives the data from the database.

## 5. **Subsequent Requests:**

The next time the same data is requested, it can be fetched directly from Memcached, reducing the load on the database.



# **Distributed System Concepts in Facebook infrastructure**


Dwarkesh

23z431

# CAP Theorem

---

- Trade-offs between Consistency, Availability, and Partition Tolerance in distributed systems.
- Facebook often chooses availability and partition tolerance over strict consistency.
- For example, in Facebook's TAO distributed database, user actions like posting, liking, or commenting are always available, even if there are network partitions.
- Consistency might be temporarily relaxed so that the system stays responsive.





# Eventual Consistency

---


- In distributed systems, data may not be immediately consistent, but it will sync eventually.
- When you 'like' a post on Facebook, it may show up instantly to you, but it might take a few seconds for that like to appear on your friend's feed.
- Facebook accepts temporary inconsistency to maintain speed and availability.
- Eventually, all replicas sync up, and everyone sees the correct, consistent state.



# Load Balancing

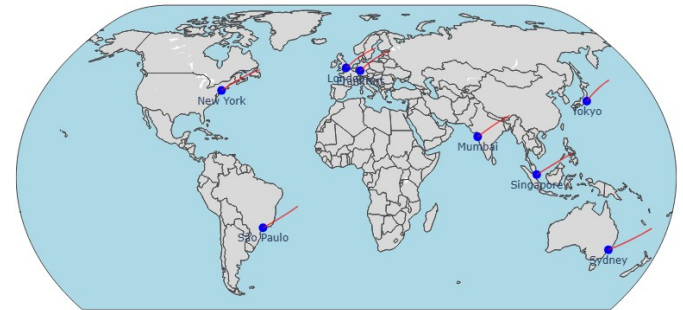
---

- Load balancing is the process of distributing incoming network traffic across multiple servers to ensure no single server is overloaded.
- Facebook's global infrastructure handles billions of user interactions per day, such as likes, comments, video plays, and messaging.
- Load balancers automatically distribute user requests to multiple backend servers to ensure fast response times and prevent crashes.



# CDN (Content Delivery Network)

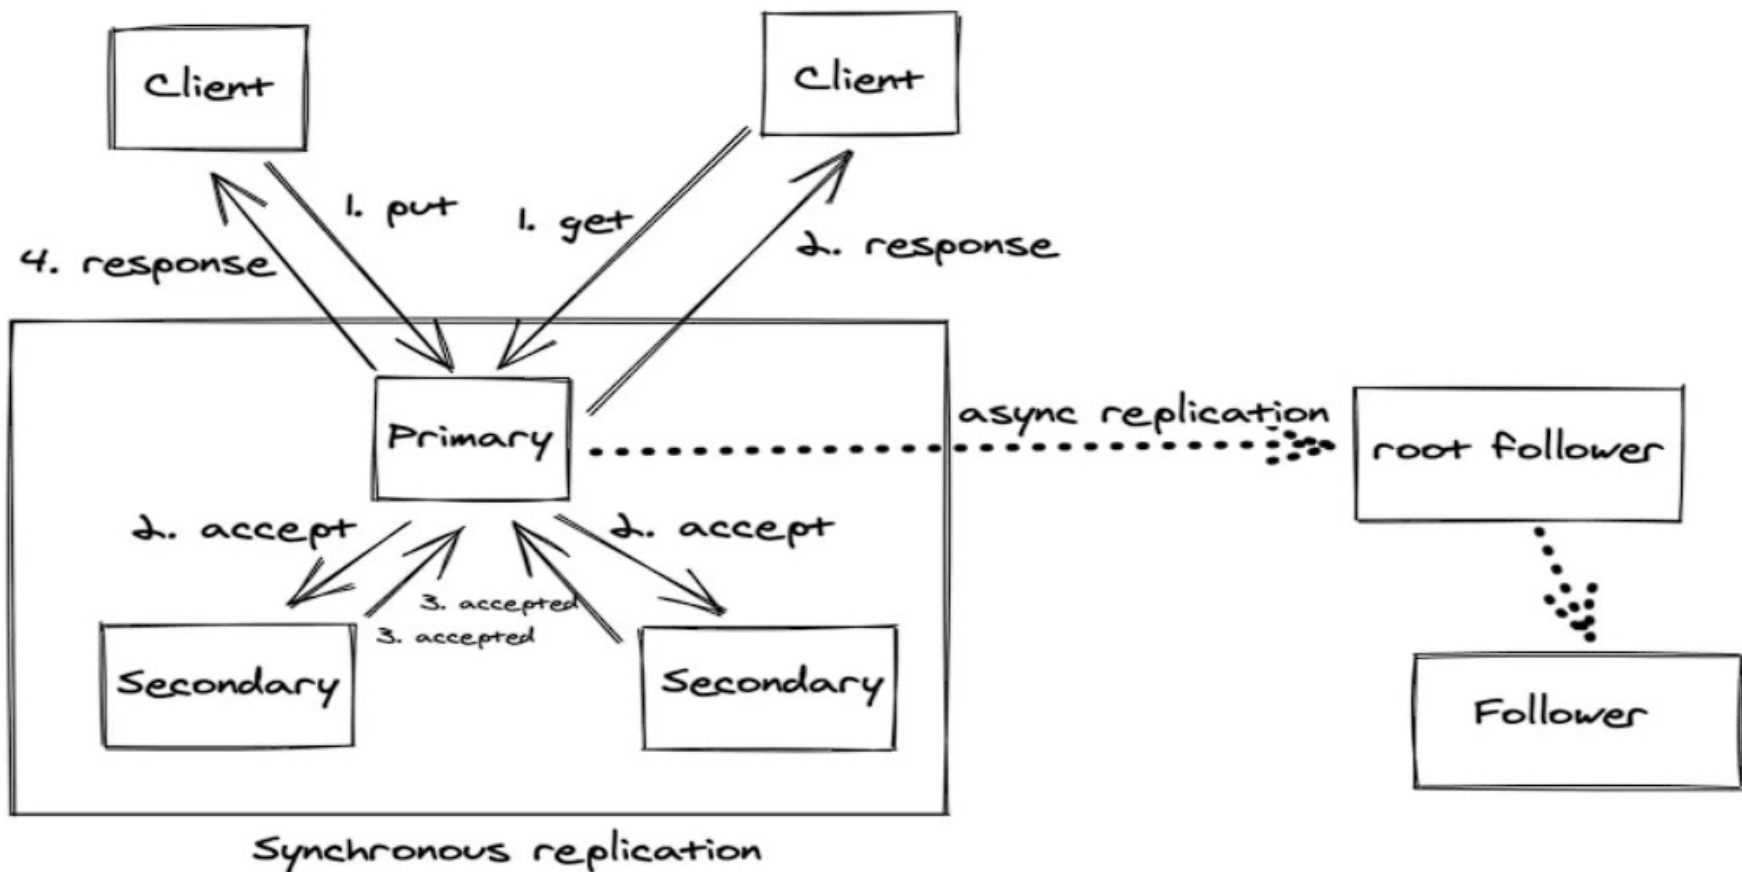
- ❖ A CDN (Content Delivery Network) is a group of distributed servers located worldwide (called edge servers) that deliver content to users from locations closer to them, reducing latency.
- ❖ Facebook's Edge Network caches images, videos, and static content (like profile pictures and stories) on edge servers near users.
- ❖ This speeds up the loading time for media-heavy content and saves bandwidth.





## **ZippyDB (Key-Value Store)**

- Developed by Facebook as a distributed key-value storage system
- Purpose: Provides low-latency data access for features like comments, notifications, and user status updates
- Architecture: Sharded design with automatic failover capabilities
- Performance: Sub-millisecond read/write operations at massive scale



## BigPipe (Dynamic Web Page Rendering)

- Problem: Traditional web pages load sequentially, creating user-perceived latency
- Solution: BigPipe breaks pages into "pagelets" that load independently and in parallel
- Implementation: Server streams content as it becomes available rather than waiting for everything

### **Benefits:**

- Significantly faster perceived page loads
- Progressive rendering of the News Feed and other content-heavy features
- Better user experience particularly on slower connections
- Legacy: Influenced modern web development frameworks and practices

## HipHop for PHP (HHVM)

**Challenge:** PHP's interpreted nature created server efficiency bottlenecks

**Solution:** HipHop transforms PHP code into highly optimized C++ then compiles it

**Evolution:**

- Started as a PHP-to-C++ translator
- Evolved into HHVM (HipHop Virtual Machine) - a JIT compiler
- Led to the development of Hack programming language

**Impact:**

- Reduced server CPU usage by up to 50%
- Enabled massive cost savings in Facebook's infrastructure
- Improved response times across the platform

---

## Scribe (Distributed Logging System)

**Purpose:** Aggregates and processes log data across Facebook's global infrastructure

**Scale:** Handles petabytes of log data daily from thousands of servers worldwide

**Architecture:**

- Distributed collection agents on every server
- Fault-tolerant message queuing
- Configurable storage backends (HDFS, S3, etc.)

**Applications:**

- Real-time monitoring and alerting
- Performance optimization
- User behavior analysis
- Security incident detection



# Thrift (Cross-Language RPC Framework)

Thrift is an **open-source Remote Procedure Call (RPC) framework** that enables **different programming languages to communicate** efficiently. It was originally developed at **Facebook** to solve the problem of **cross-language service communication**.

## Thrift solves this problem by:

- Defining a **common interface** for communication.
- Using an **efficient binary protocol** for fast data exchange.
- Supporting **multiple languages** (C++, Java, Python, PHP, Go, etc.).

# XHP (Secure HTML Rendering)

---

- XHP **prevents XSS vulnerabilities** by treating HTML as structured objects.
- It **automatically escapes user input**, making PHP-HTML integration safer.
- Used by **Facebook** to enhance security in **large-scale applications**.

If a hacker enters `<script>alert('Hacked!')</script>` as the username.


```
$userName = htmlspecialchars($_POST['name']);  
echo "<div>Hello, " . $userName . "!</div>";
```

# Autoscale (Dynamic Scaling)

---

Autoscale is a system that **automatically adjusts the number of active servers** based on **real-time traffic demands**. Instead of keeping a fixed number of servers running at all times, **autoscaling dynamically increases or decreases resources** to ensure optimal performance and cost-efficiency.

- 1.Run too many servers** all the time, wasting money.
- 2.Run too few servers**, leading to slow response times or crashes during high traffic.



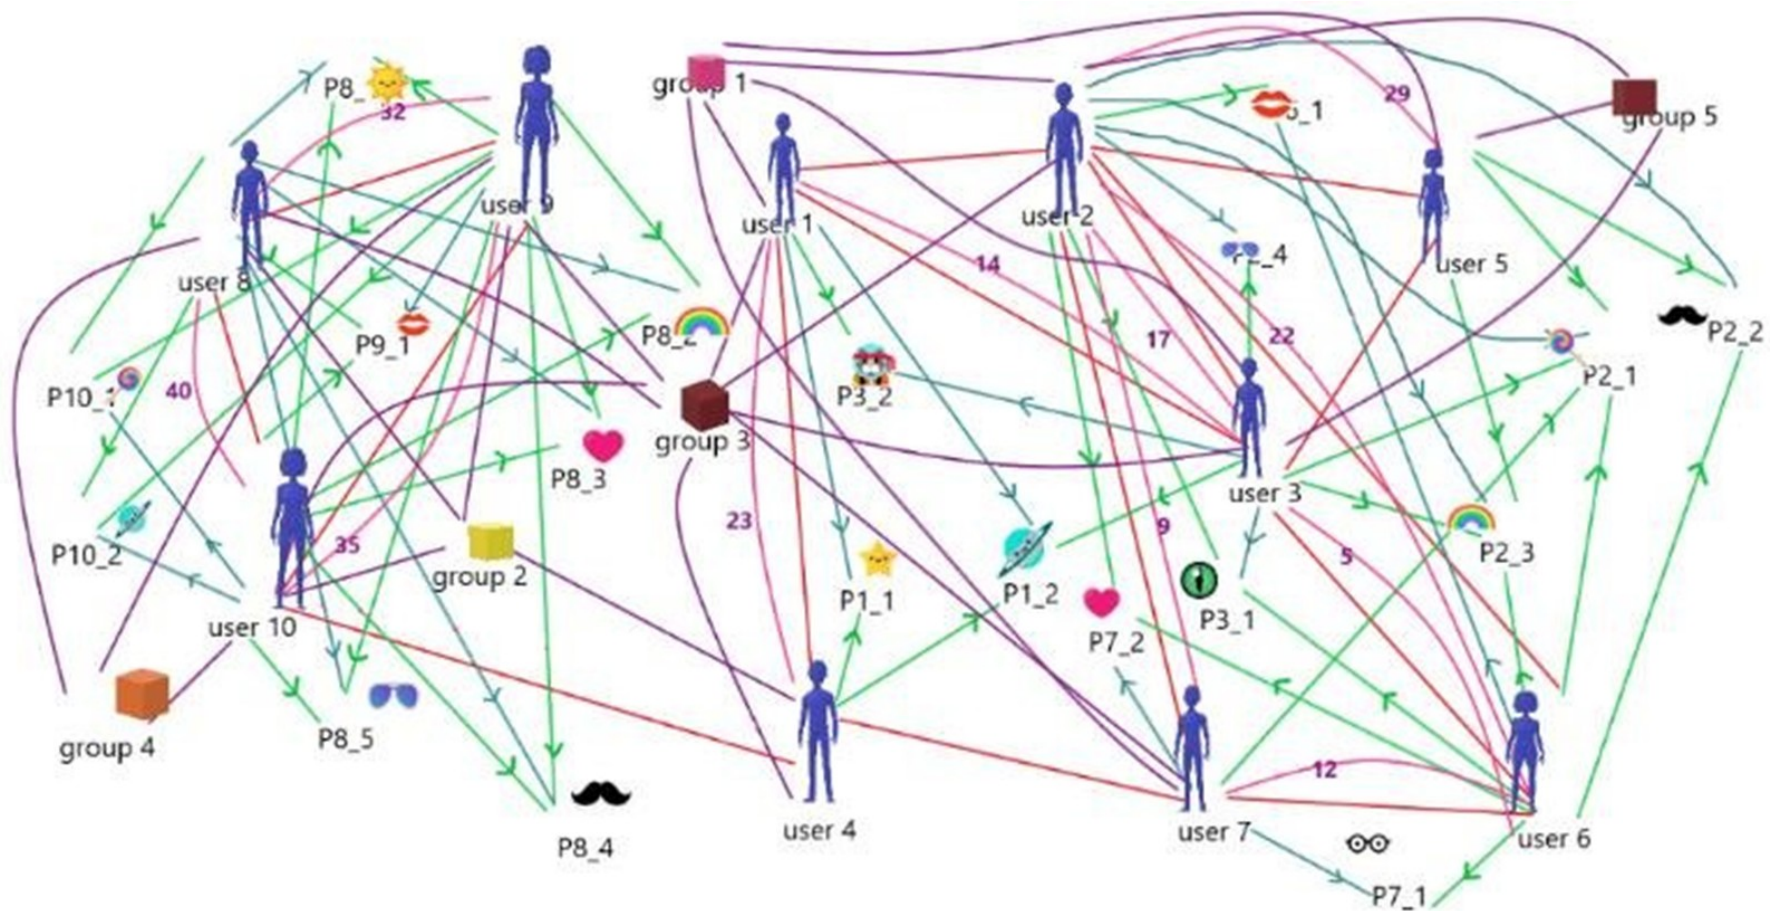
# Social Graph

---

A **social graph** is a **digital map of relationships** between people, objects, and interactions in a social network. It represents how users are connected through **friendships, likes, comments, and other interactions**.

At **Facebook**, the social graph is stored and managed using **TAO**

A social graph consists of **nodes** and **edges**.

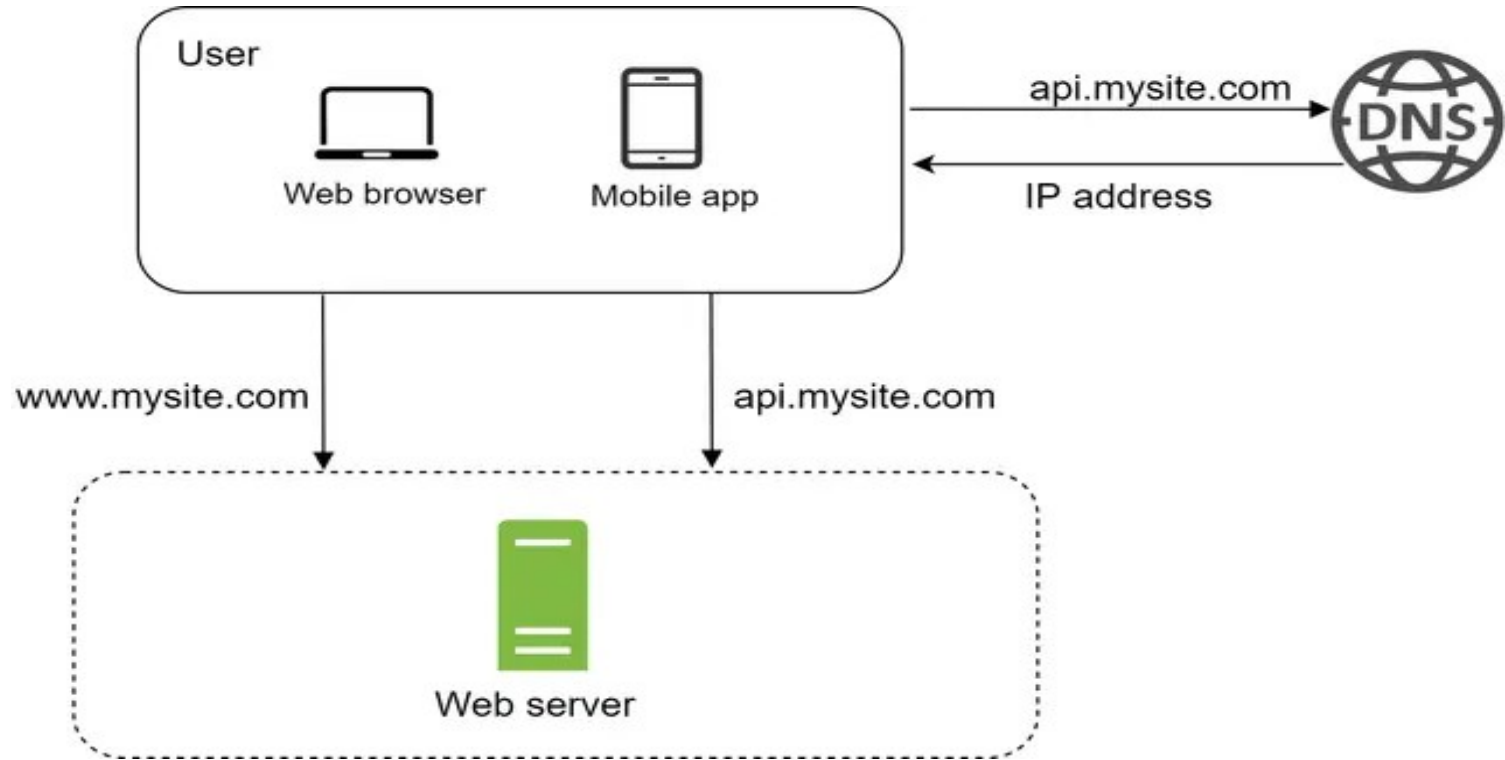


# **Scale From Zero To Millions Of Users**

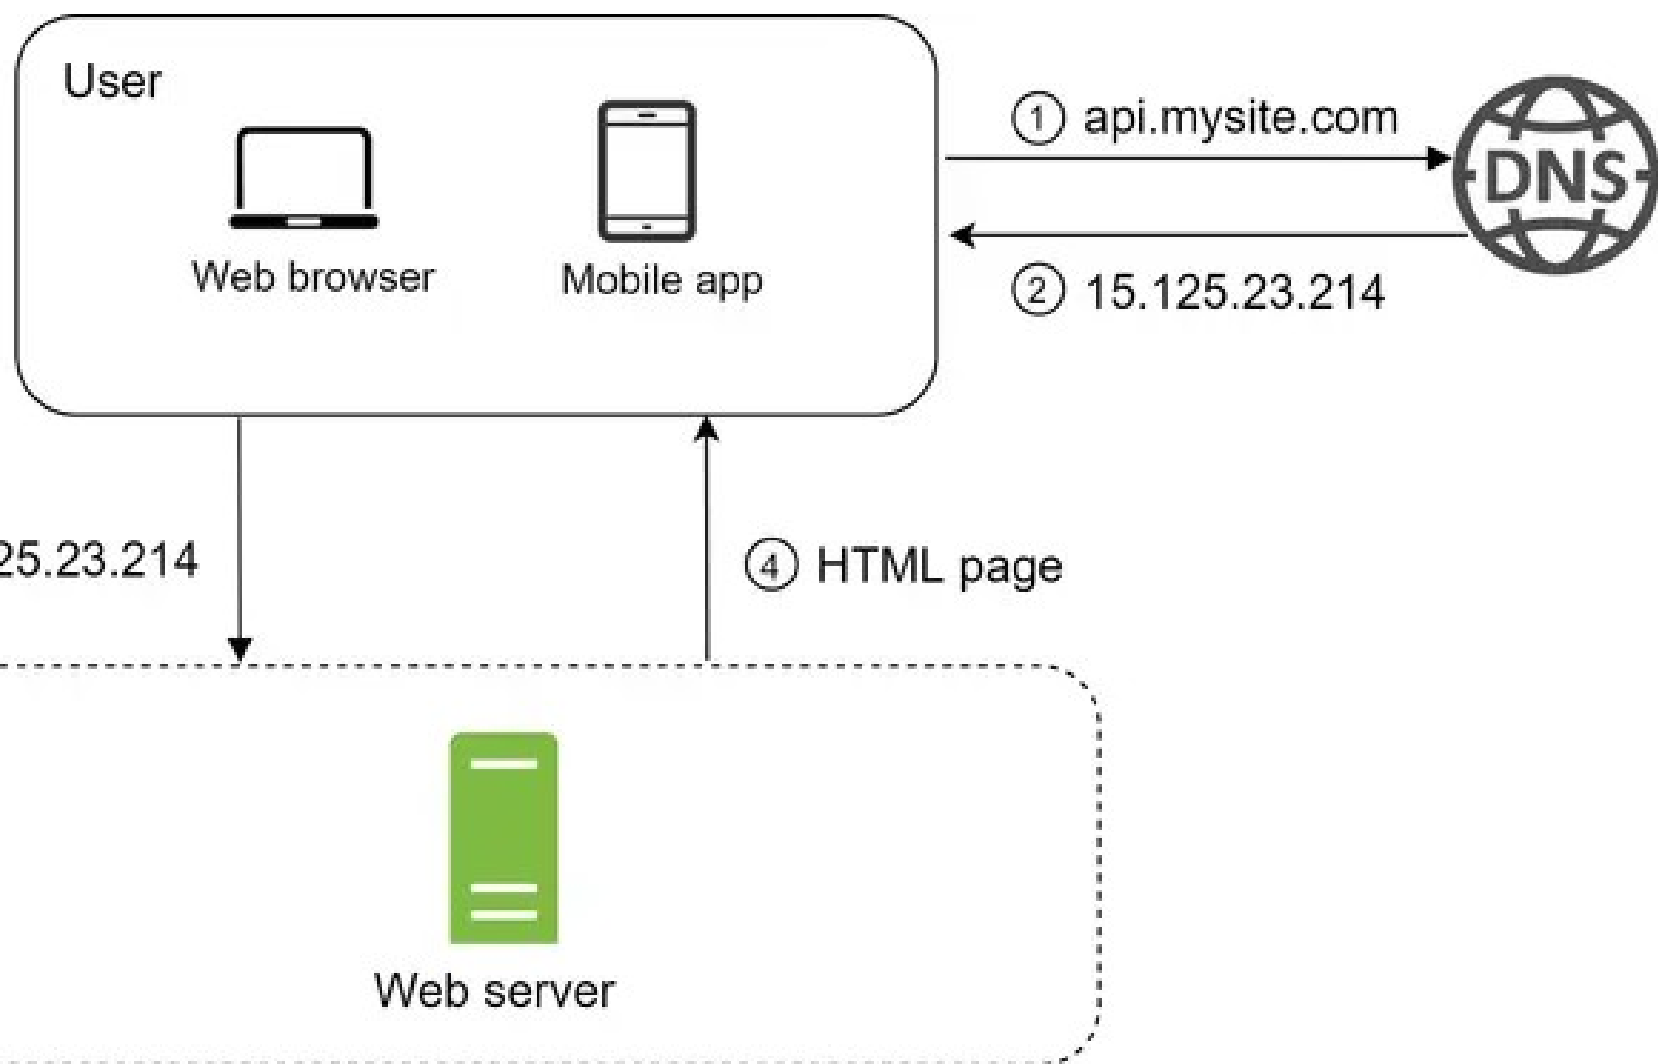
Designing a System like Facebook

Santhosh T K  
22z433

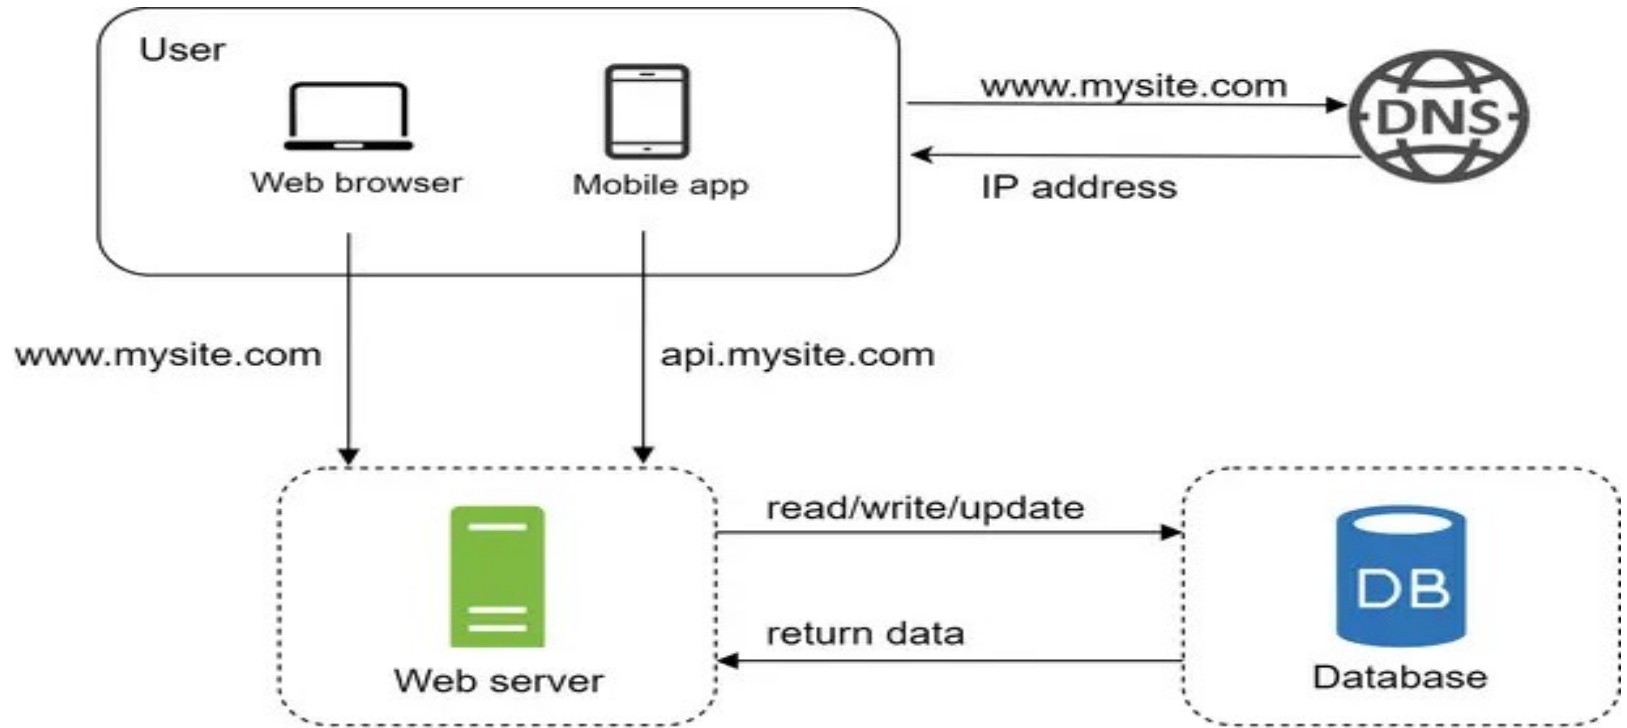
# Single server setup







# Database



# **Which DB to use Relational (or) Non-Relational ? (Design Decision)**


Go for Non-Relational if

- Need super low latency
- No relationship among data
- Need only serialization and deserialization
- Massive amount of data

If none of the above doesn't matter to you, then go with relational.


## Vertical Scaling

Increase or decrease the capacity of existing services/instances.




## Horizontal Scaling

Add more resources like virtual machines to your system to spread out the workload across them.



# Increased requests to the servers. What to do ?

Design Decision : Horizontal Scaling



Which server should handle the request ?



User



Web browser



Mobile app

Domain	IP Address
mywebsite.com	88.88.88.1

Public IP: 88.88.88.1



Load balancer

Private IP: 10.0.0.1

Private IP: 10.0.0.2

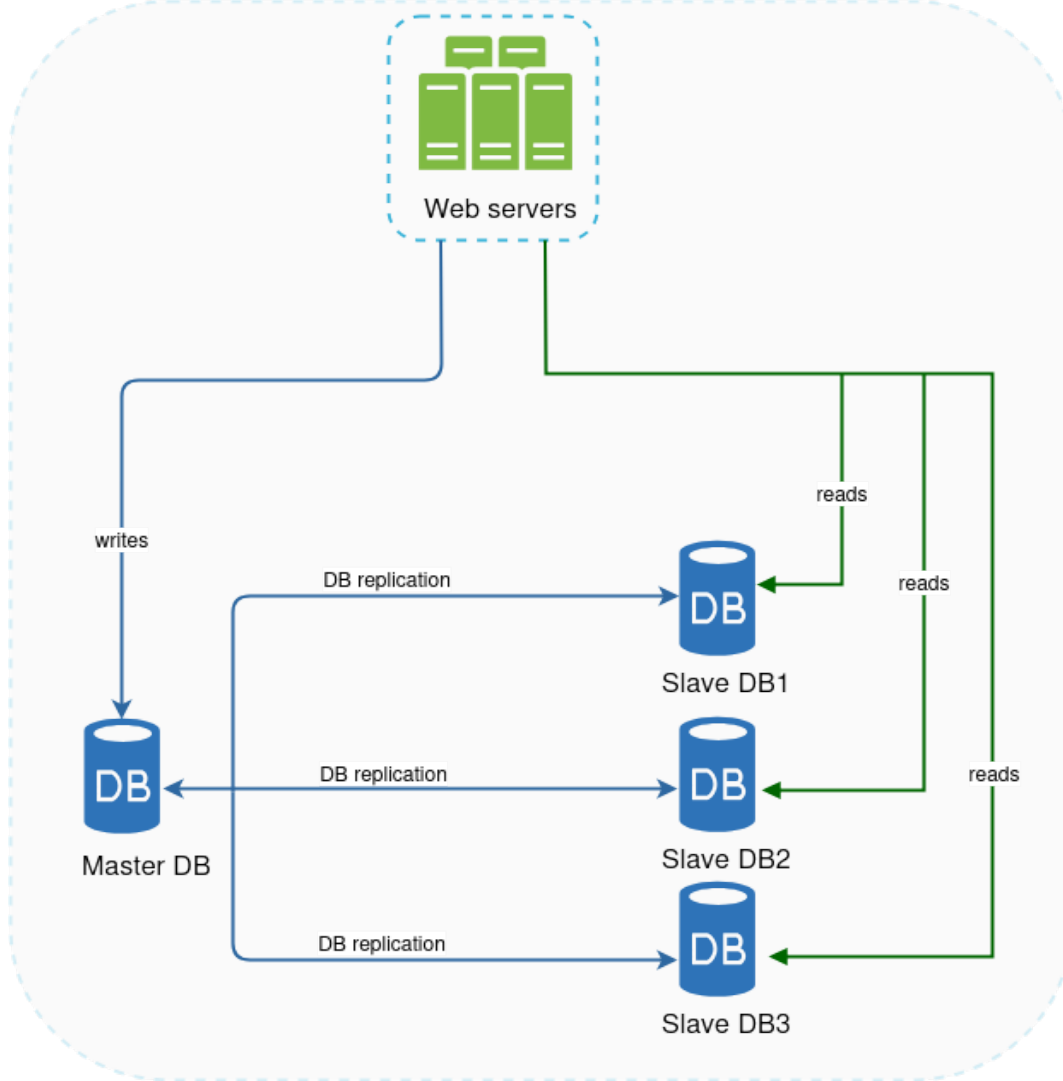


Server1

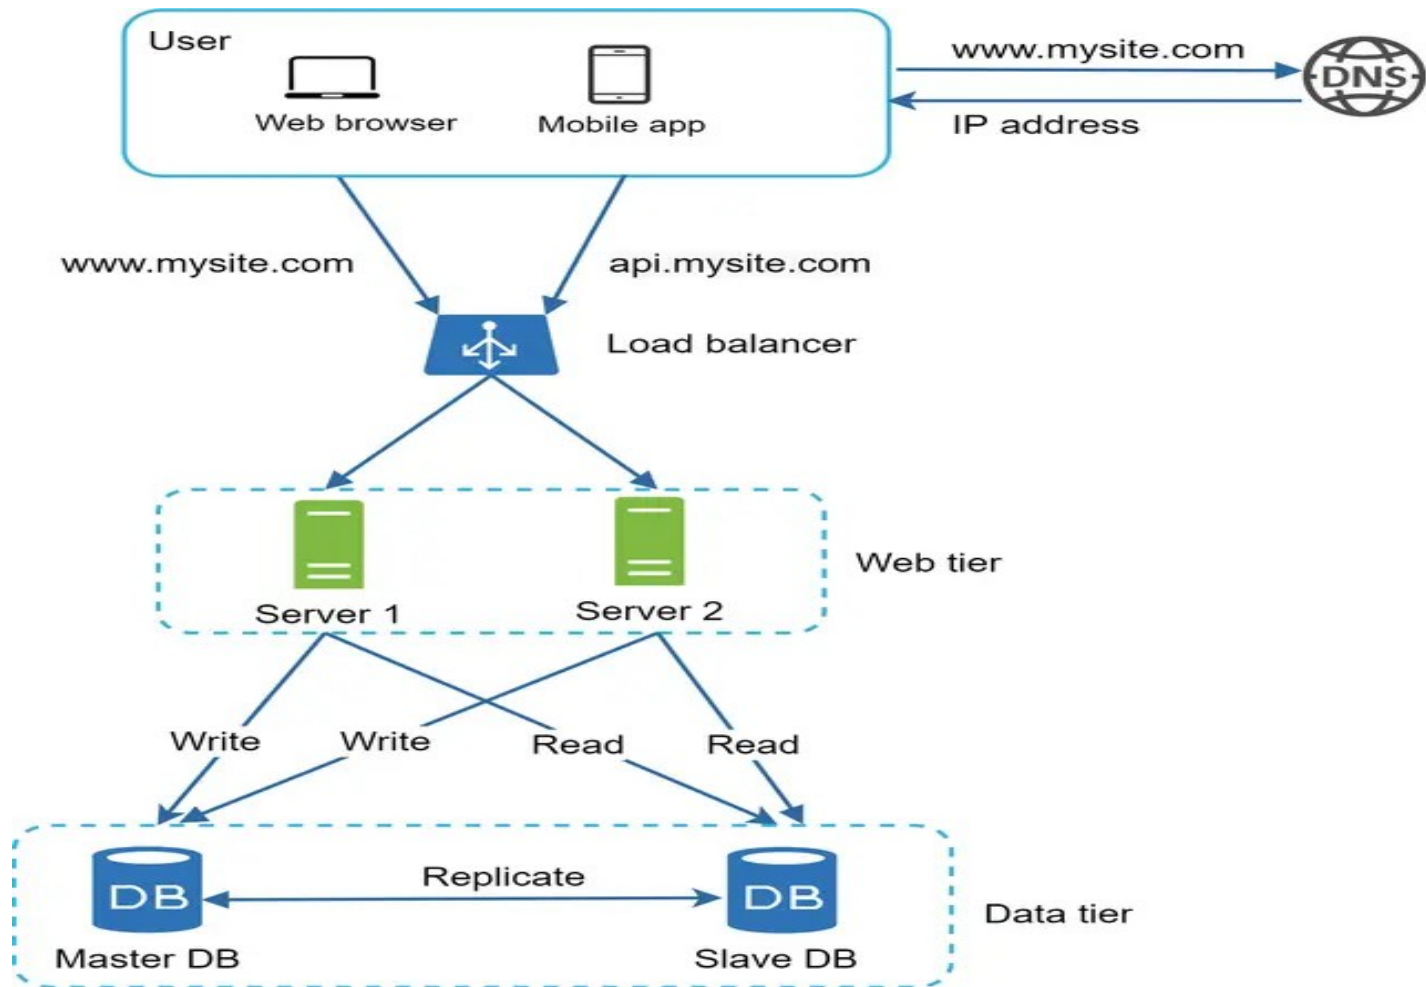


Server2

Is single database  
enough ?







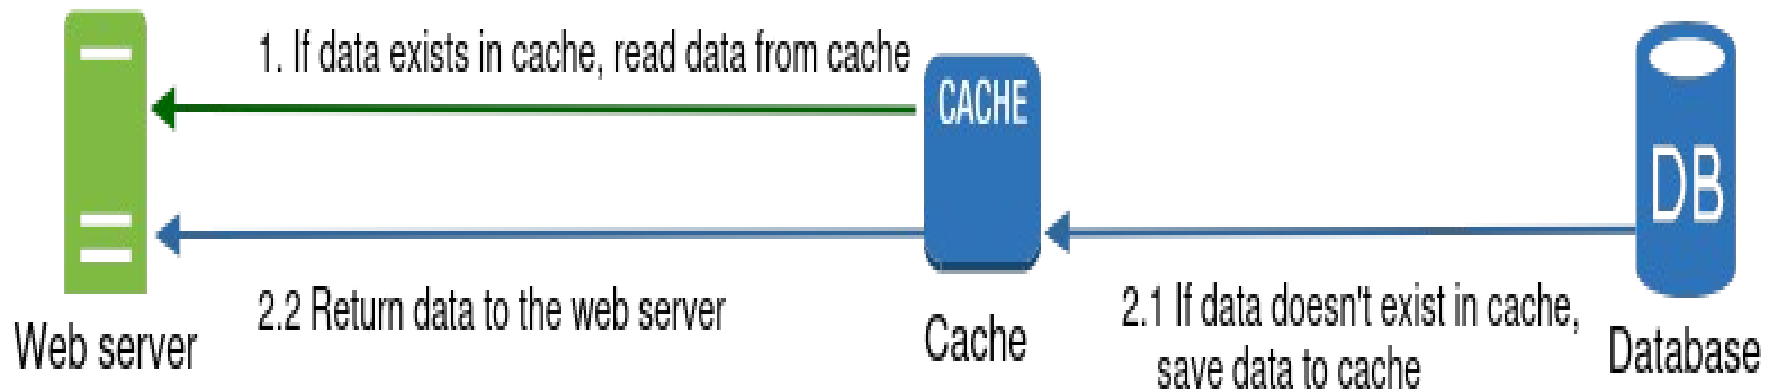
**We need to add something, Guess what ?**

Hint :


1. Trending Songs ( frequent access to the same song by different users )
2. Breaking News ( frequent access to the news )

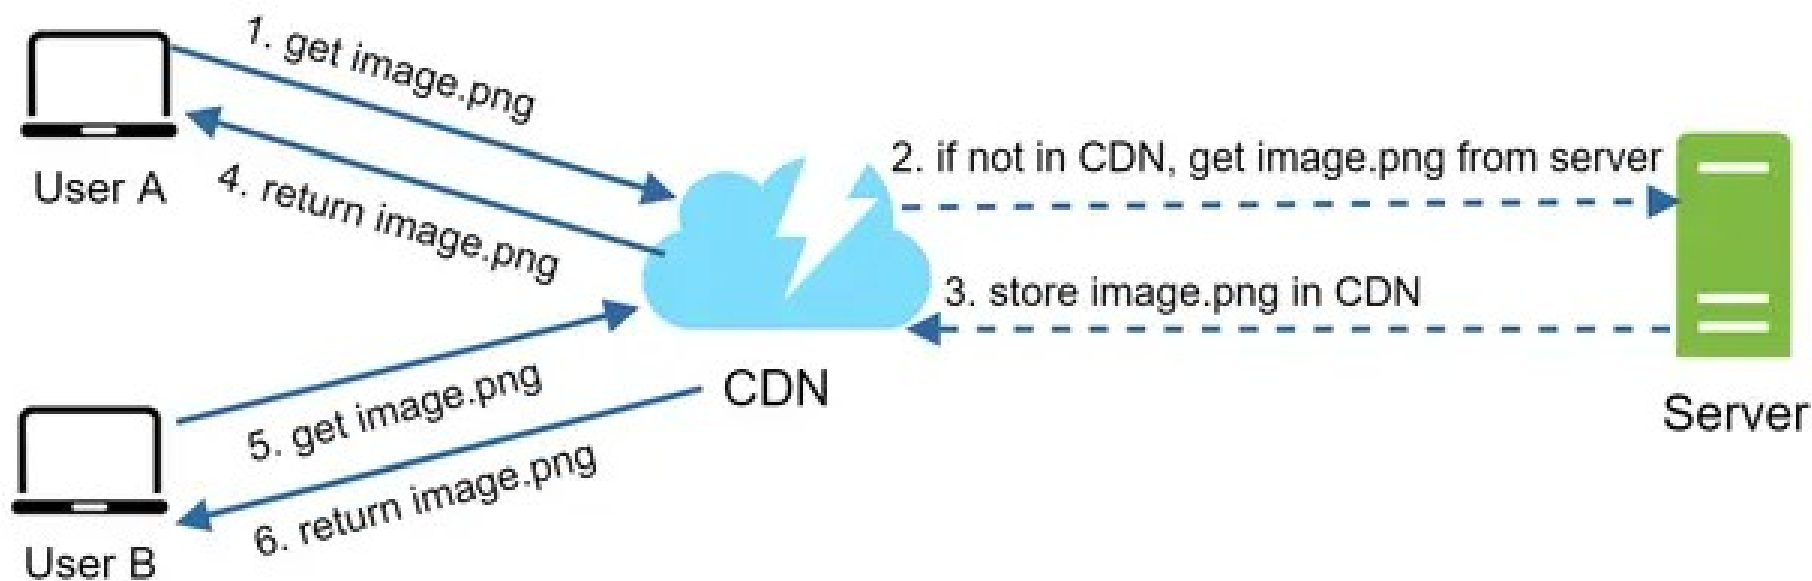
Got it right !!!

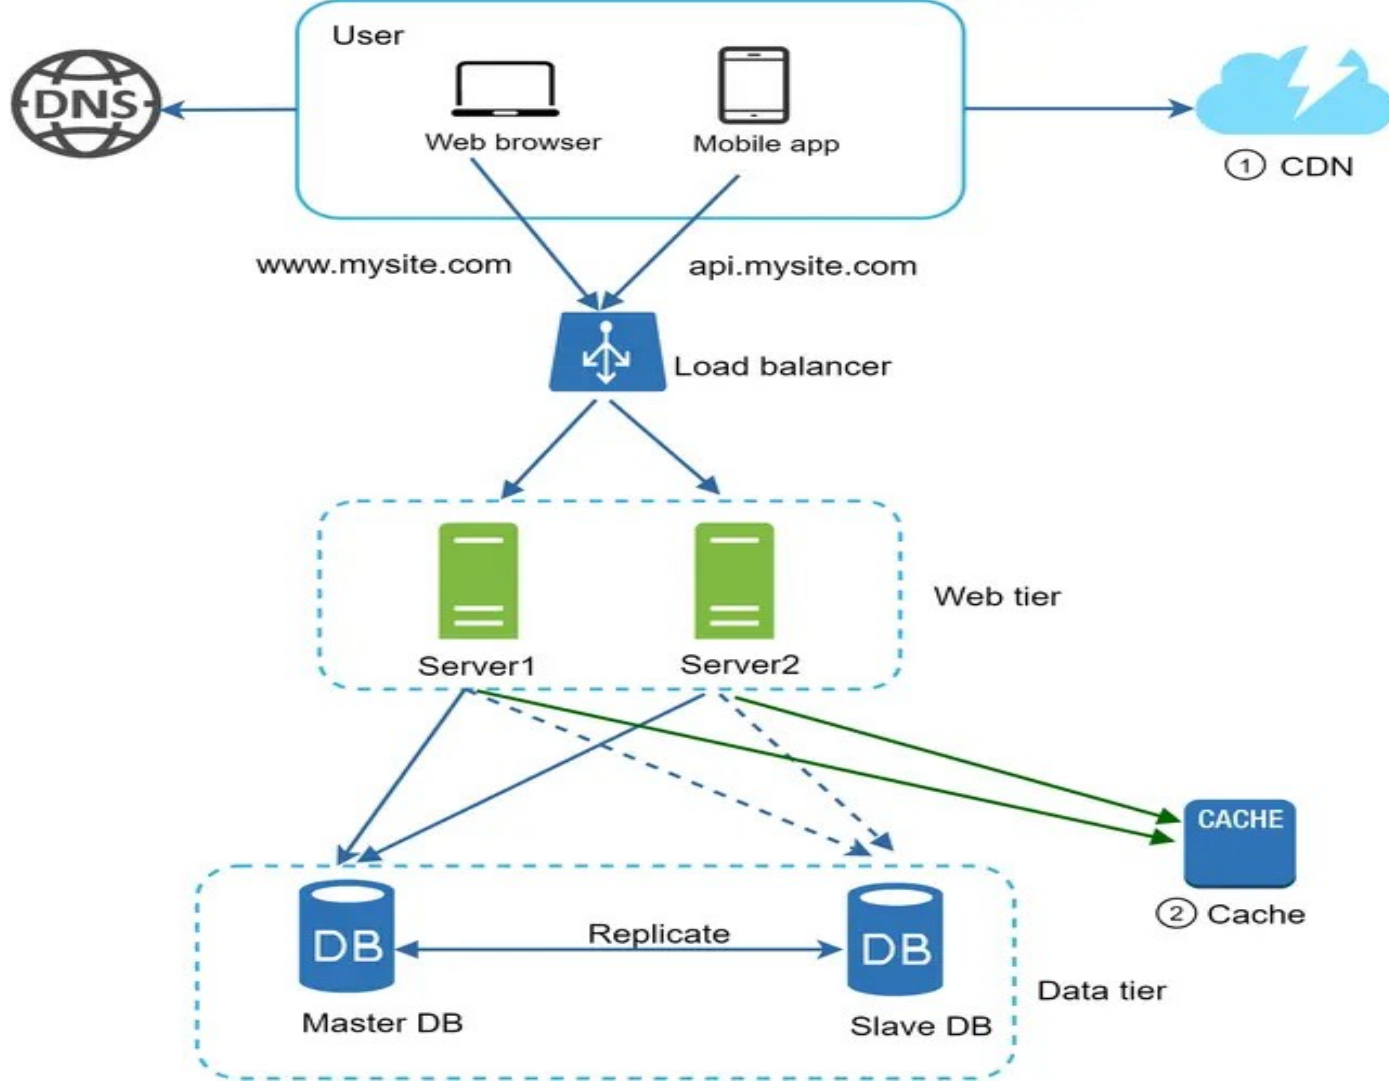
# Cache tier



# Content delivery network (CDN)

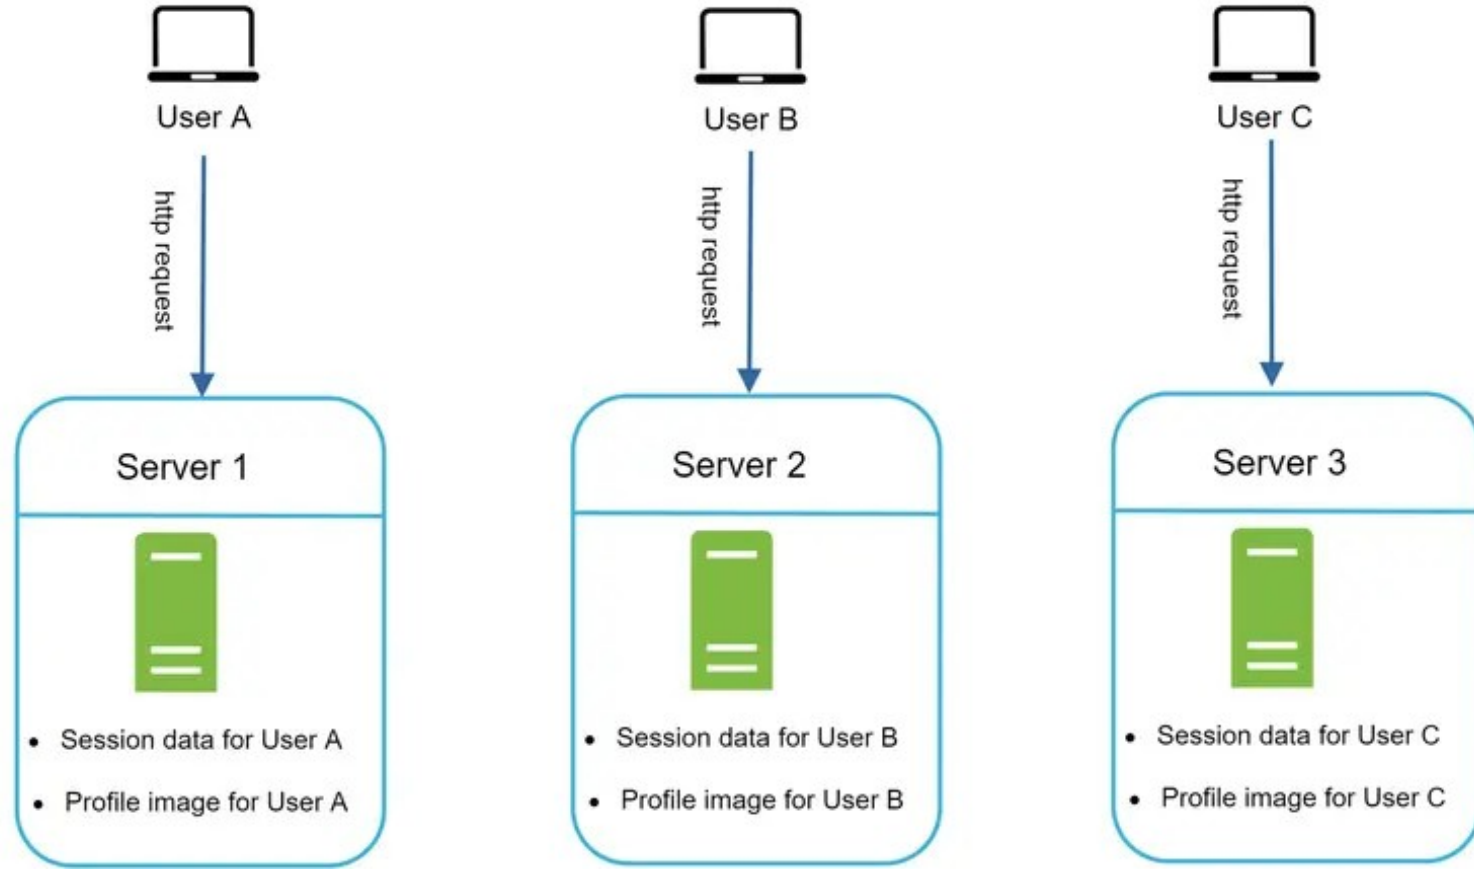







# Stateless Web Tier

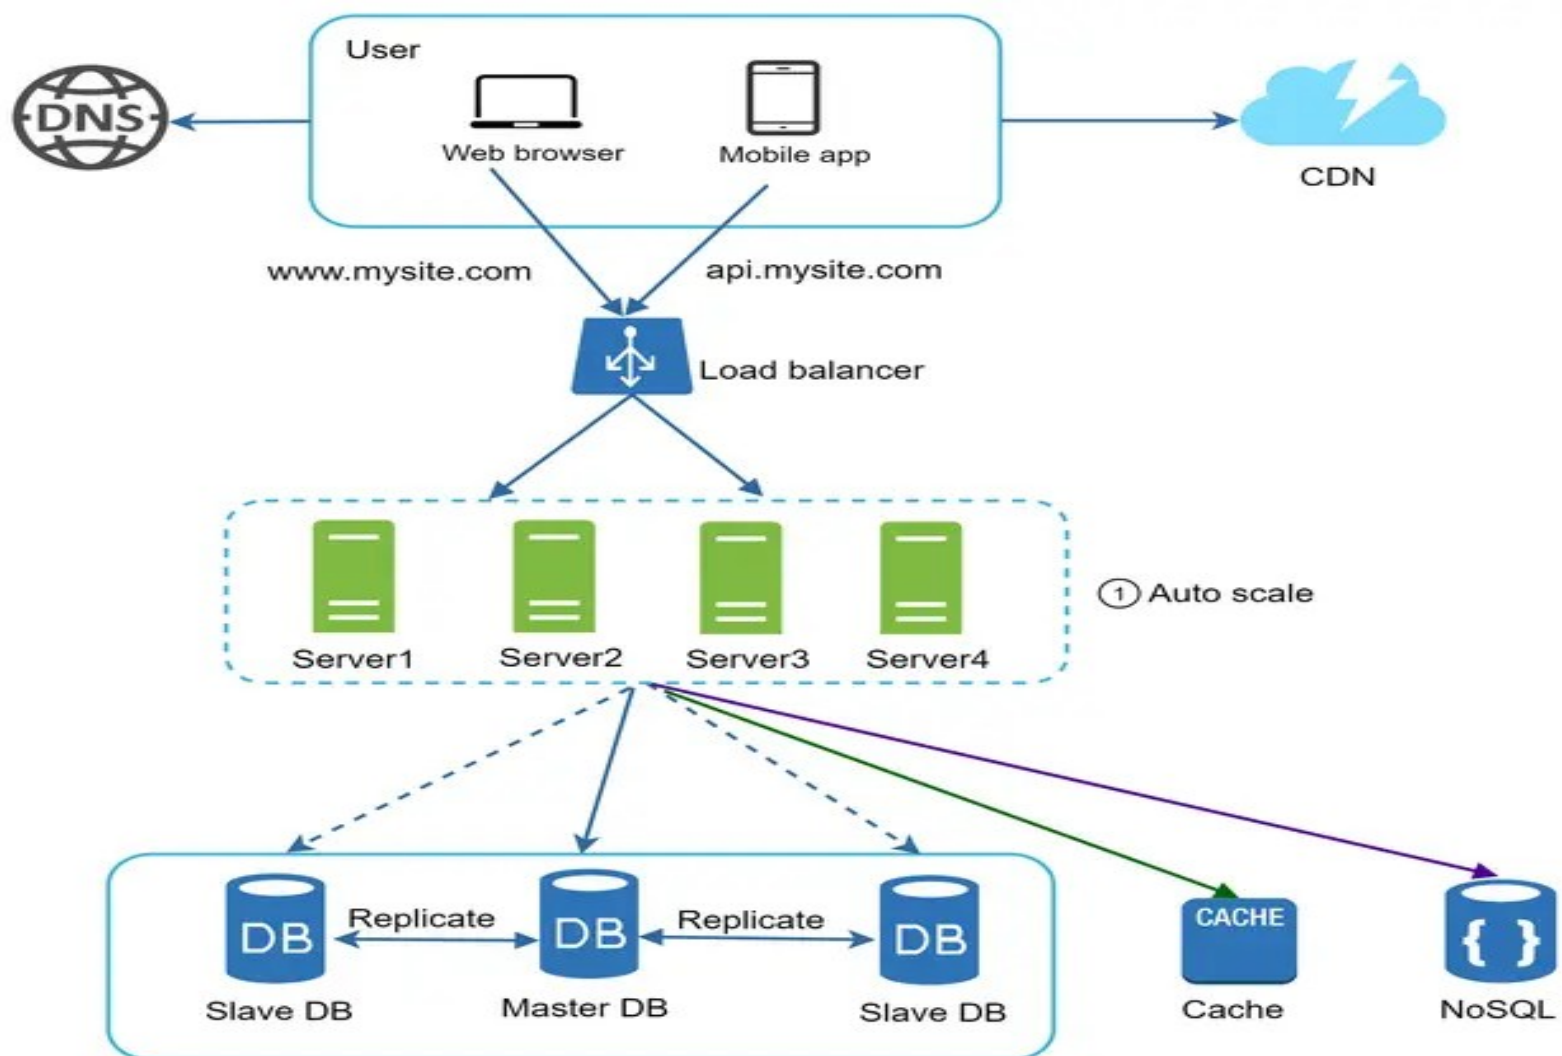
## Stateful Architecture



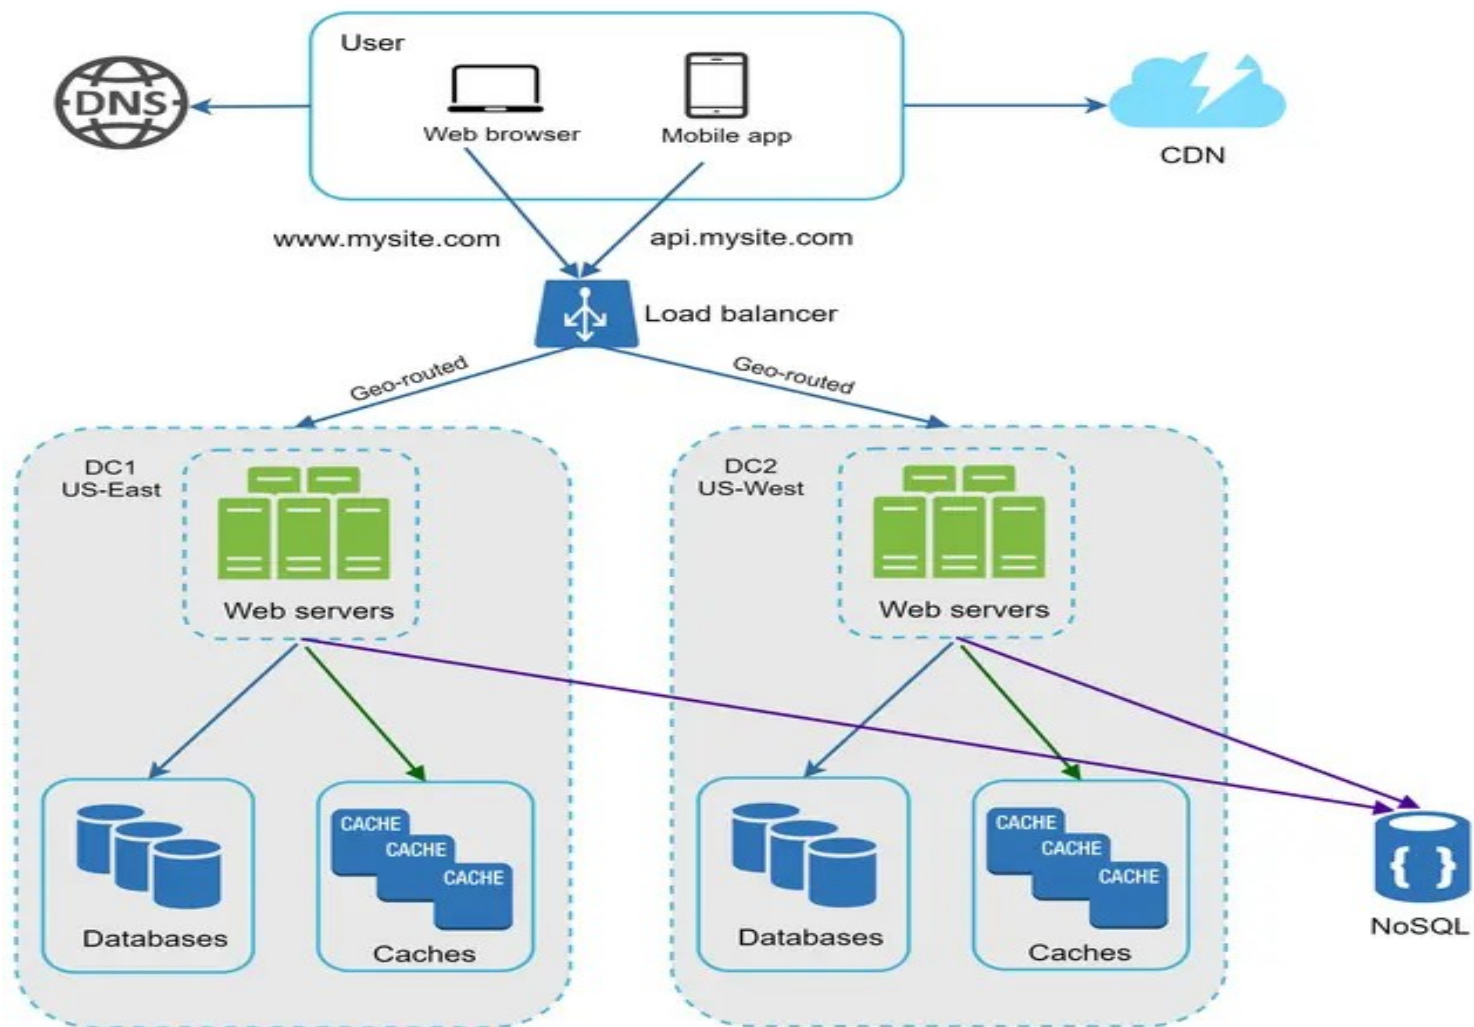


# Stateless Architecture

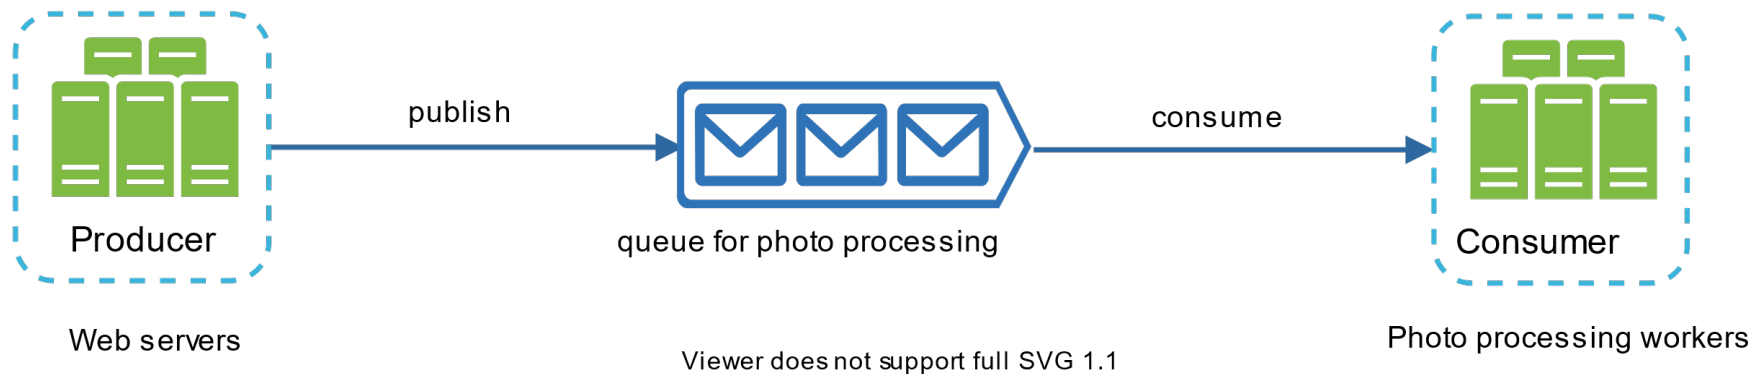


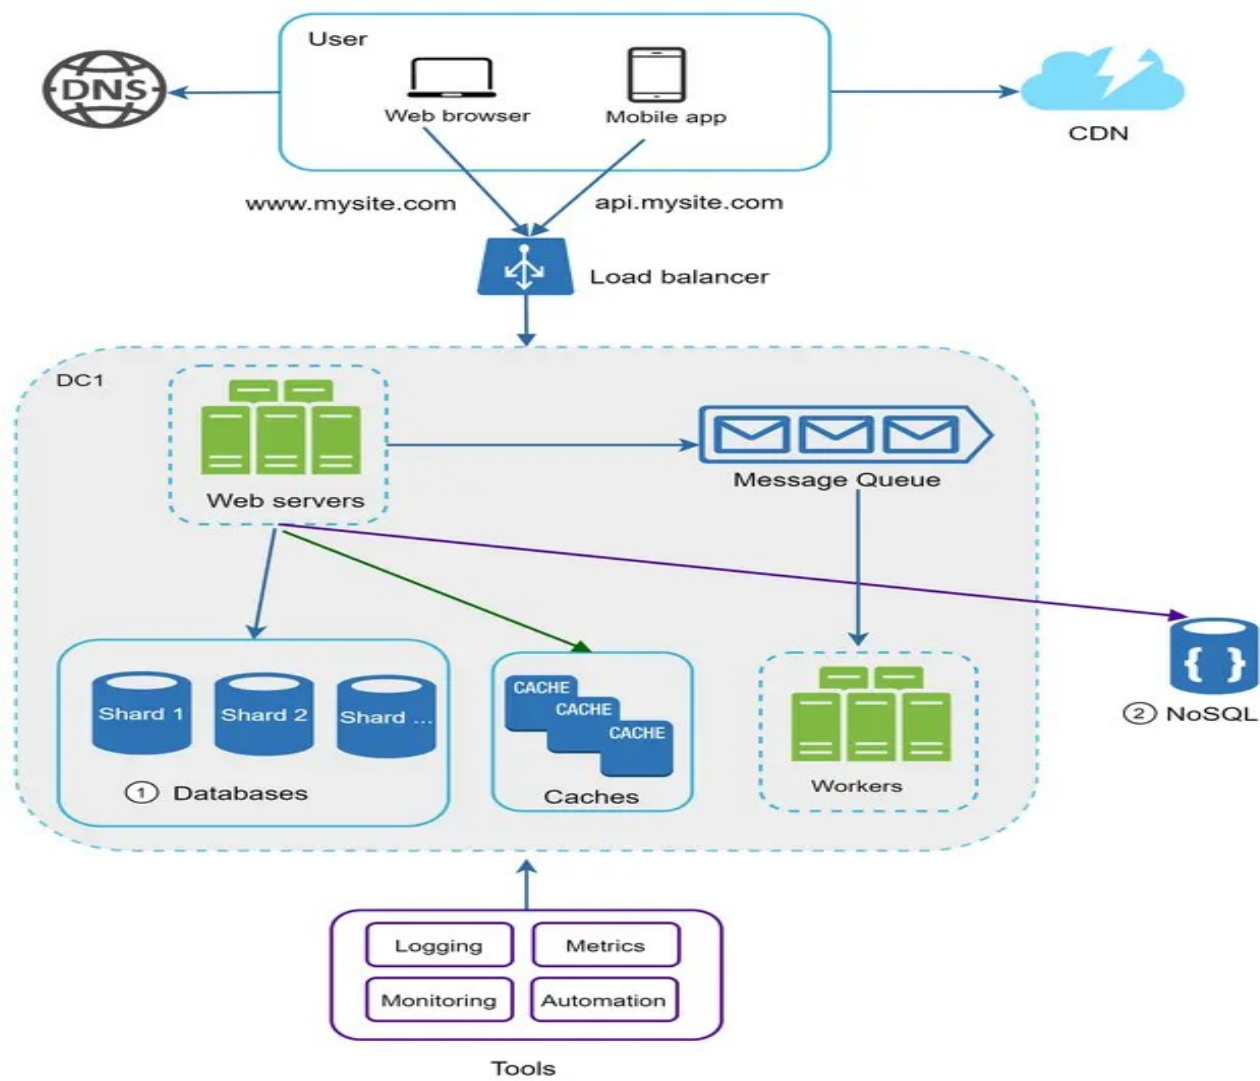


# **Data Centers**



# Message Queues for Async Communication





# References

- [1]. Scale From Zero To Millions Of Users: <https://bytebytego.com/courses/system-design-interview/scale-from-zero-to-millions-of-users>
- [2]. R. Nishtala, "Facebook, Scaling Memcache at," 10th USENIX Symposium on Networked Systems Design and Implementation (NSDI '13).
- [3]. What The Heck Are You Actually Using NoSQL For :<http://highscalability.com/blog/2010/12/6/what-the-heck-are-you-actually-using-nosql-for.html>

*Thank  
you!*