



# AMAZON EC2

ABINAYA SURESH - 22Z203

AKSHARAA P- 22Z208

HARSHINI S P - 22Z224

RAMAPRIYA S - 22Z252

RITHANIYAA B - 22Z252

THEETCHANAA RA - 22Z269





# INTRODUCTION TO AMAZON EC2

RAMAPRIYA S - 22Z252



AMAZON EC2









ON DEMAND PRICING

HOURS	672
PRICE \$	0.113
TOTAL \$	75.936







# EC2 Architecture & Key Components

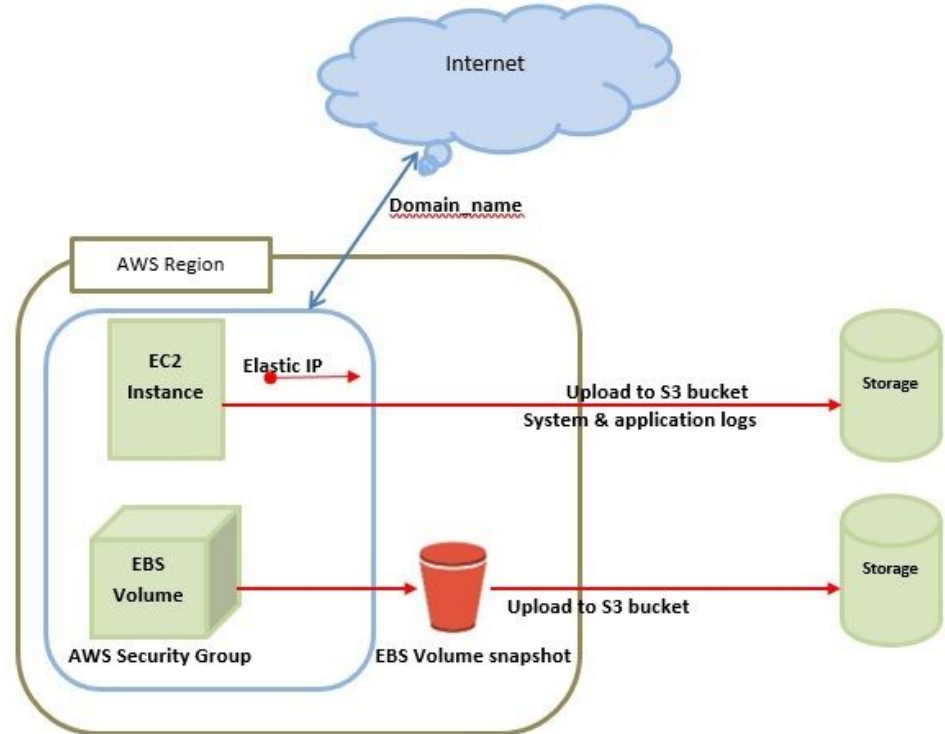
ABINAYA SURESH - 22Z203



# EC2 Architecture & Key Components

Amazon EC2 (Elastic Compute Cloud) provides scalable computing power in the AWS cloud. It is a key component of cloud-based applications, supporting high availability, security, and performance.

- Amazon Cloud-front
- Elastic Load Balancer
- Security Management
- Elastic Caches
- Elastic Block Storage
- Storage & Backups
- EC2 Instances







## Amazon CloudFront

- A Content Delivery Network (CDN) for fast delivery of dynamic, static, and streaming content.
- Uses a global network of edge locations for lower latency and improved performance.
- Works seamlessly with AWS services like S3 & EC2 and supports non-AWS origin servers.
- **Pay-as-you-go** pricing with no upfront costs.

## Elastic Load Balancer (ELB)

- Distributes incoming traffic across multiple EC2 instances for **high availability & fault tolerance**.
- Supports **auto-scaling** based on traffic fluctuations.
- Types of Load Balancers:
  - **Application Load Balancer (ALB)** – HTTP/HTTPS traffic
  - **Network Load Balancer (NLB)** – High-performance TCP/UDP traffic
  - **Classic Load Balancer (CLB)** – Basic load balancing



## Security Management in EC2

- **Security Groups:** Acts as a virtual firewall controlling inbound & outbound traffic.
- **Network ACLs:** Provides an extra layer of security at the subnet level.
- **IAM (Identity & Access Management):** Controls permissions for users and services.
- **AWS Shield:** DDoS protection for AWS applications.

## Elastic Caches

- **Amazon ElastiCache** manages in-memory caching for **fast data retrieval**.
- Reduces database load and improves **application performance**.
- Supports **Memcached** and **Redis** for caching frequently accessed data.





## Storage & Backups

- **Amazon S3:** Scalable object storage for backup & archival.
- **Amazon EBS (Elastic Block Storage):** Persistent storage for EC2, supporting snapshots & redundancy.
- **AWS Backup:** Centralized service for automated backups & compliance.

## EC2 Instances

- Virtual servers providing **scalable compute power** in AWS cloud.



# EC2 INSTANCE TYPES & USE CASES

THEETCHANAA RA - 22Z269





AWS EC2 offers different instance types optimized for various workloads.

- General Purpose
- Compute Optimized
- Memory Optimized
- Storage Optimized
- Accelerated Computing


General  
purpose




Compute  
optimized



Storage and I/O  
optimized



Memory  
optimized



GPU or FPGA  
enabled





# General-Purpose Instances



Balanced in Compute, Memory & Networking

**Best for:** Gaming servers, small databases, web apps, personal projects

## Examples:

- **T2.micro:** 1 CPU, 1 GB RAM (Free tier)
- **M6a.large:** 2 CPUs, 8GiB RAM
- **M5.large:** Intel Xeon-powered

## Features:

- AWS Graviton3 Processors - Offers better performance at a lower cost.
- EBS Optimized by default - Faster disk performance.
- Lightweight Hypervisor & Dedicated Hardware - Improves security and performance.
- Higher Bandwidth than other types

**Applications:** Web hosting, Dev/Test Environments, Content Delivery

# Compute-Optimized Instances



Optimized for High CPU Performance

**Best for:** Machine Learning, Gaming, Web servers, Batch Processing

**Examples:**

- **C5d.24 xlarge:** 96 CPUs, 192 GiB RAM, 12 Gbps Network

**Features:**

- AWS Graviton3 Processors
- DDR5 Memory (50% more bandwidth than DDR4)
- Enhanced Networking
- EBS Optimization enabled - for fast storage performance

**Applications:** Real-time multiplayer gaming, Machine Learning training

# Memory-Optimized Instances



Optimized for Large Memory Workloads

**Best for:** High-performance databases, Real-time processing

## Examples:

- **R7g.medium:** 1 CPU, 8 GiB RAM
- **X1:** 64 vCPU, 976 GiB RAM

## Features:

- Elastic Fabric Adapter (EFA) Support - for low-latency networking
- DDR5 Memory (50% more bandwidth than DDR4)
- 20% improved networking over R6g instances

**Applications:** In-memory databases, Big Data Processing

# Storage-Optimized Instances



Designed for Fast Data Access

**Best for:** Data Warehousing, High-frequency OLTP, Distributed File Systems

**Examples:**

- **Im4gn.large:** 2 CPUs, 8 GiB RAM, EBS Storage

**Features:**

- AWS Graviton2 Processors for best price-performance
- Optimized for workloads needing 4GB RAM per vCPU
- 100 Gbps Enhanced Networking (ENA) for high-throughput data processing
- Supports Local NVMe Storage (I3, D3 instances) for ultra-fast access

**Applications:** Amazon EC2 C5d/I3 instances, High-storage applications



# Accelerated Computing Instances



Uses GPUs & Hardware Accelerators for Faster Processing

**Best for:** Graphics rendering, Machine Learning, Simulations

**Examples:**

- **P4:** 8 GPUs, 96 CPUs, 1152 GiB RAM

**Features:**

- Intel Xeon Scalable Processors (3.0 GHz)
- NVIDIA A100 Tensor Core GPUs
- 400 Gbps RDMA & Elastic Fabric Adapter (EFA)

**Applications:** EC2 P3 for Deep Learning, EC2 G4 for Gaming & Video Processing



# EC2 PRICING MODELS AND COST MANAGEMENT

HARSHINI S P - 22Z224

# INTRODUCTION

- Amazon EC2 provides scalable virtual computing resources in the cloud.
- AWS offers multiple pricing models to accommodate different usage patterns and budgets.
- Users pay only for the compute resources they consume.
- **Importance of Cost Management**
  - Helps prevent unnecessary expenses by optimizing resource allocation.
  - Ensures efficient usage of cloud resources to maximize cost savings.
  - Enables scalability while maintaining control over cloud spending.



# EC2 Pricing Models



## 1. On-Demand Instances

- Pay for compute capacity per second or hour with no long-term commitment.
- Suitable for unpredictable workloads or applications with short-term needs.
- Higher cost compared to other pricing models but offers flexibility.

## 2. Reserved Instances (RI)

- Provides significant cost savings for workloads with steady-state usage.
- Requires a commitment for one or three years.
- Offers three payment options: All upfront, Partial upfront, No upfront.

## 3. Spot Instances

- Allows using unused AWS capacity at a discounted rate.
- Best for fault-tolerant workloads like batch processing and big data analysis.
- Instances can be terminated if AWS needs the capacity back.





## 4. Savings Plans

- Offers flexible discounts based on a commitment to consistent usage.
- Compute Savings Plans apply to any EC2 instance type.
- EC2 Instance Savings Plans provide better discounts but are specific to instance families.

## 5. Dedicated Hosts and Dedicated Instances

- Dedicated Hosts allow running workloads on physical servers reserved for a single customer.
- Dedicated Instances run on hardware that is not shared with other customers.
- Useful for compliance and licensing requirements.

## 6. Capacity Reservations

- Guarantees EC2 capacity in a specific Availability Zone.
- Ideal for critical applications requiring high availability.
- Charges apply even when instances are not running

# Factors affecting EC2 Costs



- **Instance Type and Size** – Costs vary based on vCPUs, memory, and configuration. Choosing the right instance optimizes expenses.
- **Region and Availability Zone** – Prices differ across AWS regions due to demand and infrastructure costs. Selecting the right region helps reduce spending.
- **Storage Costs** – EBS volumes, snapshots, and backups add to costs. Choosing the right storage type minimizes expenses.
- **Data Transfer Costs** – Inbound data is free, but outbound transfers, especially between regions or the internet, incur charges.
- **Load Balancing and Networking** – ELB and NAT Gateways increase costs. Optimizing network architecture helps control spending.

# EC2 Cost Optimization Strategies



- **Right-Sizing Instances** – Select optimal instance types based on workload needs using AWS Compute Optimizer.
- **Auto-Scaling** – Adjusts instances automatically based on demand, reducing idle costs.
- **Spot Instances** – Use for batch processing, testing, and analytics, combined with On-Demand or Reserved Instances.
- **Reserved Instances & Savings Plans** – Commit to long-term usage for cost savings on predictable workloads.
- **Monitor & Analyze Costs** – Use AWS Cost Explorer and Budgets to track spending and set alerts.
- **AWS Compute Optimizer** – Recommends instance types to avoid underutilization or over-provisioning.

# AWS Cost Management Tools



- **Cost Explorer** – Analyzes spending trends with filtering options.
- **AWS Budgets** – Sets spending limits and sends alerts when exceeded.
- **Cost & Usage Report** – Provides detailed billing insights for cost optimization.
- **Trusted Advisor** – Suggests cost-saving improvements by identifying unused resources.
- **Billing Dashboard** – Displays a breakdown of service costs and usage.

## Real-World Use Cases

- **Web Hosting** – Use On-Demand for deployment, Reserved for stable loads, and Auto-Scaling for peak traffic.
- **Big Data Processing** – Use Spot Instances, store data in S3, and process with EC2 and AWS EMR.
- **Development & Testing** – Use Spot Instances, monitor budgets, and schedule shutdowns to save costs.
- **High Availability & Disaster Recovery** – Deploy across multiple AZs, use Reserved Instances, and back up with AWS S3.





# Security, Monitoring, and Backup in EC2

RITHANIYAA B - 22Z252

# Security in Amazon EC2



## Key Security Features:

**IAM (Identity & Access Management)** – Controls who can access EC2 and what actions they can perform by defining user roles and permissions.

**Security Groups & Network ACLs** – Acts as a firewall, filtering incoming and outgoing traffic at both the instance and subnet levels to prevent unauthorized access.

**Encryption** – Protects stored data (EBS, S3, RDS) and data in transit using strong encryption protocols, ensuring information remains secure.

**AWS Shield & GuardDuty** – Shields EC2 instances from DDoS attacks and security threats by continuously monitoring and automatically mitigating risks.

# Monitoring in Amazon EC2



**Amazon CloudWatch** – Tracks **EC2 performance metrics** (CPU, memory, network activity) in real-time and sends alerts for unusual behavior.

**AWS CloudTrail** – Logs **all user activities and API actions** taken on EC2, helping with security audits and troubleshooting unauthorized access.

**AWS Trusted Advisor** – Provides **recommendations for security, cost optimization, and performance** to improve EC2 efficiency and reliability.



# Backup & Disaster Recovery in EC2



**EBS Snapshots** – Creates point-in-time backups of EC2 storage, allowing quick data restoration in case of failure or accidental deletion.

**Amazon Machine Images (AMI)** – Saves a full copy of an EC2 instance, enabling quick recovery or scaling of pre-configured setups.

**AWS Backup Service** – Automates scheduled backups of EC2 and other AWS services, ensuring data is consistently protected.

**Disaster Recovery Strategies** – Uses Multi-AZ deployment, Cross-Region Replication, and Load Balancing to maintain availability and prevent downtime.



# EC2 BENEFITS & CHALLENGES

AKSHARAA P- 2ZZ208

# BENEFITS



## **Scalability:**

- Instantly scale up during traffic spikes or business growth
- Scale down during low-demand periods to reduce costs
- Auto Scaling groups automate the process based on predefined metrics

## **Cost-effectiveness:**

- No upfront hardware investments required
- Various pricing models to optimize for your specific needs
- Detailed cost analysis and budgeting tools



### **Flexibility:**

- Over 400 instance types optimized for different workloads
- Support for Windows, Linux, and macOS
- Bare metal instances available for specialized workloads

### **Security:**

- Default encryption for data at rest and in transit
- Security groups act as virtual firewalls
- Regular security patches and updates



# CHALLENGES



## **Cost Management:**

- Unexpected costs from forgotten or idle resources
- Complex pricing models can be difficult to predict
- Data transfer costs can accumulate quickly
- Need for dedicated cost monitoring and optimization

## **Dependency on Internet Connectivity:**

- Reliance on stable internet connections
- Potential single point of failure
- Bandwidth limitations can affect performance
- Latency issues for certain applications
- Need for redundant connectivity solutions